

Cheat Sheet: LangChain Expression Language (LCEL)

Estimated time: 5 minutes

LangChain Expression Language (LCEL) is a declarative method for assembling chains from modular components within the LangChain framework. Rather than prescribing step-by-step instructions, LCEL allows you to specify the outcome, enabling LangChain to optimize its code.

This cheat sheet provides an overview of LCEL's key capabilities, patterns, and orchestration strategies.

Now let's explore some LCEL capabilities and their related benefits.

Capability	Benefits
Run optimized parallel execution	Reduces latency and increases overall application performance by running components concurrently.
Use guaranteed async support	Enables smooth, non-blocking workflows that improve responsiveness and throughput in complex chains.
Stream outputs incrementally	Delivers immediate feedback to users, lowers perceived latency, and monitors progress through each stage of the pipeline.
Trace all steps automatically with LangSmith	Provides visibility into your chain's behavior to quickly debug, monitor performance, and improve reliability.
Call all chains through a shared LCEL API	Simplifies integration and enables consistent behavior across workflows.
Deploy chains with LangServe	Accelerates the transition from development to production with minimal overhead.
Uses concise and expressive syntax	Eases connectivity among components for building robust data pipelines.

Runnables

Each component conforms to the LCEL standardized Runnable interface.

Runnables represent any component that can transform input into output. The Runnable interface provides a standard set of methods (like `invoke()`, `batch()`, `stream()`) that all components implement, making them interoperable. The following table lists some runnable types, provides a brief description of the runnable type, and lists an example use case.

Runnable type	Description	Example use case
<code>ChatModel</code>	Interfaces with LLM APIs for chat	Generate conversational responses
<code>PromptTemplate</code>	Creates formatted prompts from variables	Prepare structured inputs for LLMs
<code>OutputParser</code>	Converts raw outputs to structured data	Extract structured data from LLM responses
<code>RunnableLambda</code>	Wraps custom Python functions	Implement custom business logic
<code>RunnableSequence</code>	Chains multiple Runnables together	Create multi-step processing pipelines

Next, learn about using runnable chains.

Runnable chains: Use cases, components, and workflows.

You can build runnable chains by connecting components from LangChain's standardized Runnable interface. Each component passes its output directly to the next component, creating a workflow.

Next, review some runnable chain use cases, their components, and their workflow summaries.

Use cases	Components	Workflow summary
Simple question answering	<code>PromptTemplate</code> → <code>ChatModel</code> → <code>StrOutputParser</code>	Format a question, send to LLM, return text response
Retrieval augmented generation (RAG)	<code>Retriever</code> → <code>PromptTemplate</code> → <code>ChatModel</code> → <code>OutputParser</code>	Find relevant documents, combine with prompt, generate response
Function calling	<code>PromptTemplate</code> → <code>ChatModel</code> → <code>Tool</code>	Format prompt, generate function call, parse function parameters, execute function
Structured output	<code>PromptTemplate</code> → <code>ChatModel</code> → <code>JsonOutputParser</code>	Format prompt, generate response, parse to JSON object

LCEL function reference

LCEL provides a rich set of functions to invoke, compose, and configure Runnables. This reference section categorizes LCEL functions into basic operations, composition techniques, data manipulation, advanced patterns, configuration tools, and streaming/batching support. Each function enhances how you structure and control your chains. It's important to note that functions that begin with the letter `a` are asynchronous functions.

Basic operations

Function	Description	Usage
----------	-------------	-------

<code>invoke()</code> / <code>ainvoke()</code>	Execute a Runnable with a single input	Process one input and get one output
<code>batch()</code> / <code>abatch()</code>	Process multiple inputs efficiently in parallel	Run the same operation on multiple inputs at once
<code>stream()</code> / <code>astream()</code>	Return incremental results as they're generated	Show partial responses as they're created

Composition patterns

Function	Description	Usage
Pipe operator <code> </code> or <code>.pipe()</code>	Create a sequence of Runnables	Chain components where output of one becomes input to the next
<code>RunnableParallel</code>	Execute multiple Runnables with the same input, concurrently	Process the same input in different ways simultaneously
<code>RunnableLambda</code>	Convert Python functions into Runnables	Add custom logic within a chain

Data manipulation patterns

Function	Description	Usage
<code>RunnablePassthrough.assign()</code>	Add new fields to the input dictionary	Augment input with additional data while preserving the original input
<code>RunnablePassthrough()</code>	Return input unchanged	Include the original input as part of the output
<code>.pick()</code>	Select specific keys from the dictionary output	Filter output to only needed fields

Advanced patterns

Function	Description	Usage
<code>.bind()</code>	Set default values for parameters	Fix certain parameters while leaving others configurable
<code>.with_fallbacks()</code>	Try alternative Runnables if the primary fails	Handle errors by providing backup components
<code>.with_retry()</code>	Add automatic retry capability	Retry operations on failure For example: Network issues

Configuration

Function	Description	Usage
<code>configparameter</code>	Control runtime execution	Pass to <code>invoke()</code> , <code>batch()</code> , <code>stream()</code> methods with settings like concurrency limits and tracing parameters
<code>.with_config()</code>	Create Runnable with default configuration	Apply the same configuration to all invocations automatically

Streaming and batching

Function	Description	Usage
<code>astream_events()</code>	Detailed stream of execution events	Monitor the entire execution process, including intermediate steps
<code>batch_as_completed()</code> / <code>abatch_as_completed()</code>	Process inputs in parallel, return as completed	Start processing results as soon as they're available

Next, explore some recommended approaches to orchestration.

Recommended approaches to orchestration

When building LLM applications, you can choose the orchestration method based on the complexity of your use case. The following table provides some guidance for when to use a direct call, LCEL, or LangGraph.

Use case	When to use	Recommended tool
Single LLM call	- You need to generate text from a prompt - The overhead of setting up a chain is not justified	LLM directly
Simple chains	- You need a straightforward pipeline of components. For example: A prompt + LLM + parser + basic retrieval - Your application can benefit from parallelized code and streaming - Your application has a clear linear flow with minimal branching	LCEL

Complex logic with branching	• Your application requires complex state management • Your application requires conditional flows, loops, or cycles • You are implementing multi-agent systems that interact with each other	LangGraph
------------------------------	---	-----------

Author(s)

IBM Skills Network Team



Skills Network